



Figure 1: Extended results for the static experiment with varying ratios of data used to train the downstream classifier. The remaining data is used for evaluation and the underlying embedding remains identical.